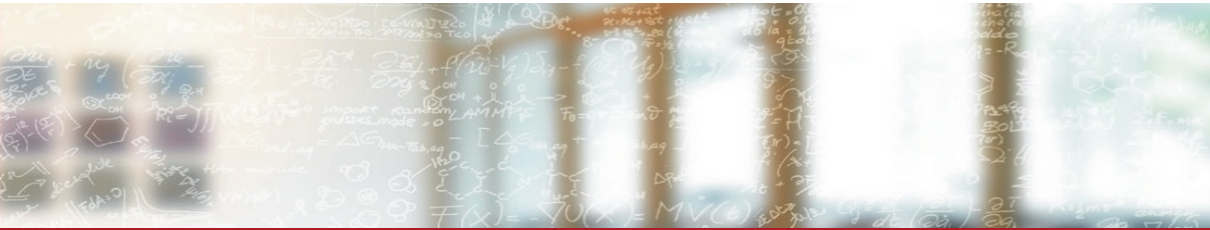




**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Introduction to OpenACC

Directive Based GPU Programming

*Vasileios Karakasis, CSCS*

May 14–15, 2018

# Overview

- Day 1 (Understand the basic concepts of OpenACC)
  - Execution and memory model
  - Basic directives
  - Profiling and debugging
  - Hands-on sessions
- Day 2 (Advanced topics)
  - Asynchronous execution and wait queues
  - Interoperability with CUDA and MPI
  - Deep copy
  - Hands-on sessions

# What is OpenACC?

- Collection of compiler directives for specifying loops and regions to be offloaded from a host CPU to an attached accelerator device
- Host + Accelerator programming model
- High-level representation
- Current specification version: 2.6

# Why to use OpenACC?

Because ...

- I don't care about all the little hardware details, I want my science done.
- I want to run on accelerators, but I still need a fast and readable code.
- I need portability across different accelerator vendors, but also be able to run on the multicore.
- I inherited a large legacy monolithic codebase, which I don't dare to refactor completely, but I need to get my results faster.
- My code is in Fortran.

# OpenACC is not a silver bullet

- A high-level representation is not a panacea.
  - You still need to understand and adapt to the programming model.
- Does not substitute hand-tuning, but can serve as a very good starting point.
- User base not yet as large as of classic OpenMP for multicores, but it is expanding.
  - You may run into compiler bugs or specification ambiguities.

# Format of directives

- C/C++
  - `#pragma acc directive-name [clause-list] new-line`
  - Scope is the following *block of code*
- Fortran
  - `!$acc directive-name [clause-list] new-line`
  - Scope is until `!$acc end directive-name`

# Programming model

- Host-directed execution
- Compute intensive regions are offloaded to attached accelerator devices
- Host orchestrates the execution on the device
  - Allocations on the device
  - Data transfers
  - Kernel launches
  - Wait for events
  - Etc. . .

# Execution model

- The device executes *parallel* or *kernel regions*
- Parallel region
  - Work-sharing loops
- Kernel region
  - Multiple loops to be executed as multiple kernels
- Levels of parallelism
  1. *Gang*
  2. *Worker*
  3. *Vector*
  - Parallelism levels are decided by the compiler but can be fine-tuned by the user



# Execution model

- The device executes *parallel* or *kernel regions*
- Parallel region
  - Work-sharing loops
- Kernel region
  - Multiple loops to be executed as multiple kernels
- Levels of parallelism
  1. *Gang* → *CUDA block*
  2. *Worker* → *CUDA warp*
  3. *Vector* → *CUDA threads*
  - Parallelism levels are decided by the compiler but can be fine-tuned by the user
  - Mapping to CUDA blocks/warps/threads is implementation defined

# Execution model



## Modes of execution

- Gang
  - Gang-redundant (GR)
  - Gang-partitioned (GP)
- Worker
  - Worker-single (WS)
  - Worker-partitioned (WP)
- Vector
  - Vector-single (VS)
  - Vector-partitioned (VP)

# Execution model

## Modes of execution



```
for (int i = 0; i < 4096; ++i) {  
    a[i] = c*b[i];  
}
```



# Execution model

## Modes of execution


```
for (int i = 0; i < 4096; ++i) {  
    a[i] = c*b[i];  
}
```



# Execution model

## Modes of execution

```
for (int i = 0; i < 4096; ++i) {  
    a[i] = c*b[i];  
}
```



```
for (int i = 0; i < 1K; ++i) {  
    a[i] = c*b[i];  
}  
for (int i = 1K; i < 2K; ++i) {  
    a[i] = c*b[i];  
}  
for (int i = 2K; i < 3K; ++i) {  
    a[i] = c*b[i];  
}  
for (int i = 3K; i < 4K; ++i) {  
    a[i] = c*b[i];  
}
```

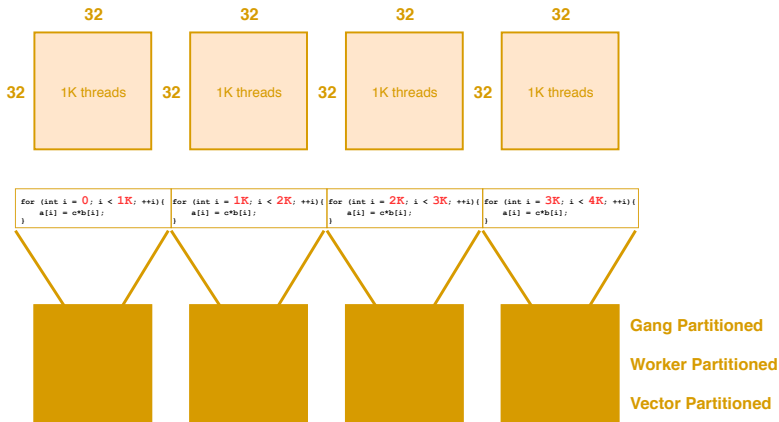


Gang Partitioned  
Worker Partitioned  
Vector Single

# Execution model

## Modes of execution

```
for (int i = 0; i < 4096; ++i) {  
    a[i] = c*b[i];  
}
```



# Execution model

## The kernels construct

### Multiple loops inside kernels construct

```
!$acc kernels
  !GR mode
  do i = 1, N
    !compiler decides on the partitioning (GP/WP/VP modes)
    y(i) = y(i) + a*x(i)
  enddo
  do i = 1, N
    !compiler decides on the partitioning (GP/WP/VP modes)
    y(i) = b*y(i) + a*x(i)
  enddo
!$acc end kernels
```

- Compiler will try to deduce parallelism
- Loops are launched as different GPU kernels

# Execution model

The `parallel` construct

## Parallel construct

```
!$acc parallel
  do i = 1, N
    ! loop executed in GR mode
    y(i) = y(i) + a*x(i)
  enddo
!$acc loop
do i = 1, N
  !compiler decides on the partitioning (GP/WP/VP modes)
  y(i) = b*y(i) + a*x(i)
enddo
!$acc end parallel
```

- No automatic parallelism deduction → parallel loops must be specified explicitly
- Implicit gang barrier at the end of `parallel`



# Execution model

## Work-sharing loops

- C/C++: `#pragma acc loop`
  - Applies to the immediately following `for` loop
- Fortran: `!$acc loop`
  - Applies to the immediately following `do` loop
- Loop will be automatically striped and assigned to different threads
  - Use the `independent` clause to force striping
- Convenience syntax combines `parallel`/kernels and loop constructs
  - `#pragma acc parallel loop`
  - `#pragma acc kernels loop`
  - `!$acc parallel loop`
  - `!$acc kernels loop`

# Execution model

## Work-sharing loops – the collapse clause

### Collapse loops

```
!$acc loop collapse(2)
do i = 1,N
    do j = 1,N
        A(i,j) = coeff*B(i,j)
    enddo
enddo
```

#### ■ OpenACC vs. OpenMP

- OpenACC: apply the `loop` directive to the following  $N$  loops and possibly collapse their iteration spaces if independent
- OpenMP: Collapse the iteration spaces of the following  $N$  loops

# Execution model

## Controlling parallelism

- Amount of parallelism at the kernels and **parallel** level
  - `num_gangs(...), num_workers(...), vector_length(...)`
- At the loop level
  - `gang, worker, vector`

100 thread blocks with 128 threads each

```
!$acc parallel num_gangs(100), vector_length(128)
  !$acc loop gang, vector
    do i = 1, n
      y(i) = y(i) + a*x(i)
    enddo
!$acc end parallel
```

# Execution model

## Variable scoping

- Allowed in the `parallel` directive only
- By default, if outside of a code block, variables are shared in global memory
- `private`: A copy of the variable is placed in each *gang* (CUDA block)
- `firstprivate`: Same as `private` but initialized from the host value

# Execution model

## Variable scoping

- Allowed in the `parallel` directive only
- By default, if outside of a code block, variables are shared in global memory
- `private`: A copy of the variable is placed in each *gang* (CUDA block)
- `firstprivate`: Same as `private` but initialized from the host value

## Implicit scoping:

- (C/C++/Fortran) Loop variables are private to the *thread* that executes the loop
- (C/C++ only) Scope of variables declared inside a parallel block depends on the current execution mode:
  - *Vector-partitioned* mode → private to the thread
  - *Worker-partitioned, Vector-single* mode → private to the worker
  - *Worker-single* mode → private to the gang

# Execution model

## Reduction operations

- `#pragma acc parallel reduction(<op>:<var>)`
  - e.g., `#pragma acc parallel reduction(+:sum)`
- `#pragma acc loop reduction(<op>:<var>)`
- `var` must be scalar
- `var` is copied and default initialized within each gang
- Intermediate results from each gang are combined and made available outside the parallel region
- Complex numbers are also supported
- Operators: `+`, `*`, `max`, `min`, `&`, `|`, `%`, `&&`, `||`

# Execution model

Calling functions from parallel regions

- `#pragma acc routine {gang | worker | vector | seq}`
  - Just before the function declaration or definition
- `!$acc routine {gang | worker | vector | seq}`
  - In the specification part of the subroutine
- Parallelism level of the routine
  - `gang`: must be called from GR context
  - `worker`: must be called from WS context
  - `vector`: must be called from VS context
  - `seq`: must be called from sequential context

# Memory model

Where is my data?

- The host and the device have separate address spaces
  - Data management between the host and the device is the programmer's responsibility
  - You must make sure that all the necessary data for a computation is available on the accelerator before entering the compute region
  - You must make sure to transfer the processed data back to the host if needed




# Memory model

Where is my data?

- The host and the device have separate address spaces
  - Data management between the host and the device is the programmer's responsibility
  - You must make sure that all the necessary data for a computation is available on the accelerator before entering the compute region
  - You must make sure to transfer the processed data back to the host if needed
- But there can be some exceptions:
  - The “device” might be the multicore → no need for data management
  - Some compilers may infer automatically the necessary data transfers
  - Nvidia Pascal GPUs provide efficient support for a unified memory view between the host and the accelerator


# Memory model

Separate address spaces




# Memory model

Separate address spaces




# Memory model

Separate address spaces




# Memory model

Separate address spaces




# Memory model

Separate address spaces




# Memory model

Separate address spaces



# Memory model


Separate address spaces





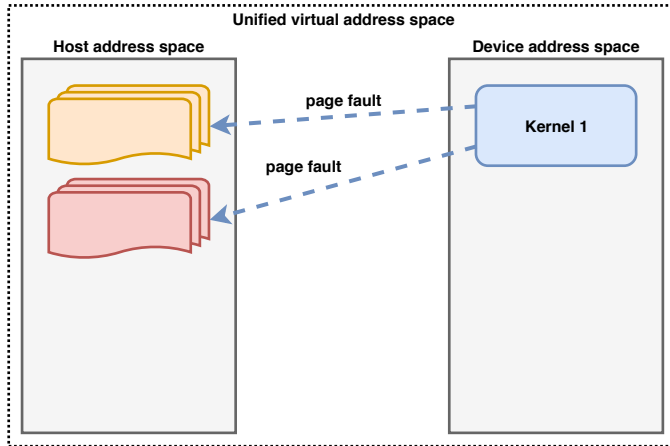
# Memory model

Unified memory address space




# Memory model

Unified memory address space



# Memory model

Unified memory address space



# Memory model

Directives accepting data clauses

Data clauses may appear in the following directives:

- Compute directives:
  - `#pragma acc kernels`
  - `#pragma acc parallel`
- Data directives:
  - `#pragma acc data`
  - `#pragma acc enter data`
  - `#pragma acc exit data`
  - `#pragma acc declare`
  - `#pragma acc update`

# Memory model

## Data clauses

- `create(a[0:n])`: Allocate array `a` on device
- `copyin(a[0:n])`: Copy array `a` to device
- `copyout(a[0:n])`: Copy array `a` from device
- `copy(a[0:n])`: Copy array `a` to and from device
- `present(a)`: Inform OpenACC runtime that array `a` is on device
- `delete(a)`: Deallocate array `a` from device (exit data only)

Not for the `acc update` directive

# Memory model

## The `acc data` directive

- Defines a scoped data region
  - Data will be copied in at entry of the region and copied out at exit
  - A *structural reference count* is associated with each memory region that appears in the data clauses
- C/C++: `#pragma acc data [data clauses]`
  - The next block of code is a data region
- Fortran: `!$acc data [data clauses]`
  - Defines a data region until `!$acc end data` is encountered

# Memory model

The `acc enter/exit data` directives

- Defines an unscoped data region
  - Data will be resident on the device until a corresponding `exit data` directive is found
  - Useful for managing data on the device across compilation units
  - A *dynamic reference count* is associated with each memory region that appears in the data clauses
- C/C++:
  - `#pragma acc enter data [data clauses]`
  - `#pragma acc exit data [data clauses]`
- Fortran:
  - `!$acc enter data [data clauses]`
  - `!$acc exit data [data clauses]`

# Memory model

The `acc declare` directive

- Functions, subroutines and programs define *implicit data regions*
- The `acc declare` directive is used in variable declarations for making them available on the device during the lifetime of the implicit data region
- Useful for copying global variables to the device
  
- C/C++: `#pragma acc declare [data clauses]`
- Fortran: `!$acc declare [data clauses]`



# Memory model

## The `acc update` directive

- May be used during the lifetime of device data for updating the copies on either host or the device
- `#pragma acc update host(<var-list>)`
  - Update host copy with corresponding data from the device
- `#pragma acc update device(<var-list>)`
  - Update device copy with corresponding data from the host

# Memory model

## Array ranges

Data clauses may accept as arguments

- Whole arrays

- C/C++: You *must* specify bounds for dynamically allocated arrays
  - `#pragma acc data copyin(a[0:n])`
  - But `#pragma acc data present(a)` is acceptable: a's bounds can be inferred by the runtime
- Fortran: array shape information is already embedded in the data type
  - `!$acc data copyin(a)`

- Array subranges

- `#pragma acc data copyin(a[2:n-2])`

# Synchronization directives

- Atomic operations
  - `#pragma acc atomic [atomic-clause]`
  - `!$acc atomic [atomic-clause]`
  - Atomic clauses: read, write, update and capture
  - Example of “capturing” a value:
    - `v = x++;`
- No global barriers → cannot be implemented due to hardware restrictions
- No equivalent of `__syncthreads()`

# Leverage the unified memory

- Virtual address space shared between CPU and GPU
- The CUDA driver and the hardware take care of the page migration
- Introduced with the Kepler architecture and CUDA 6, but is significantly improved with Pascal

# Leverage the unified memory

- Virtual address space shared between CPU and GPU
  - The CUDA driver and the hardware take care of the page migration
  - Introduced with the Kepler architecture and CUDA 6, but is significantly improved with Pascal
- 
- You could completely omit the data management in OpenACC !
  - Supported by the PGI compiler using the `-ta=tesla:managed` option

# Hands-on exercises

## General information

The initial course material is available on Github and it will be update during the course:

- `git clone https://github.com/vkarak/openacc-training.git`
- `git pull origin master` to get the latest version

Directory structure:

- `practicals/`: The hands-on exercises
- `scripts/`: Set up scripts to make your life easier
- `slides/`: Slides of the course
- `ci/`: Continuous integration tests for the exercises (ask me offline if interested)

# Hands-on exercises

## General information

- `grep TODO *.{cpp,f90,f03}`
- Both Cray/PGI compilers are supported, unless otherwise stated
  - Suggest using PGI for the advanced examples
- `source <repodir>/scripts/setup_pgi.sh` → will make available PGI 18.4
- `module load craype-accel-nvidia60` for loading CUDA and set the target architecture to the GPU
- `make`
  - For Cray compiler you may use `make VERBOSE=1` to get diagnostics information

# Hands-on

## Exercise 1 – AXPY

- `practicals/axpy/axpy_openacc.{cpp,f90}`
- Run as:  
`srun --reserv=openacc -Cgpu ./axpy.openacc [ARRAY_SIZE]`
  - `ARRAY_SIZE` is power of 2, default is 16
- Try with different sizes. Does the GPU outperform the CPU version?



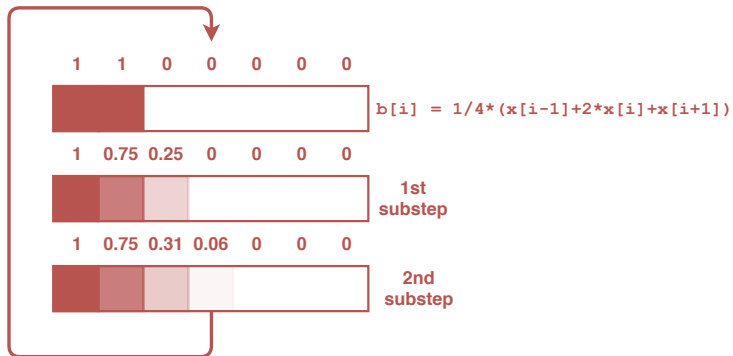
# Hands-on

## Exercise 2 – Dot product

- `practicals/basics/dot_openacc.{cpp,f90}`
- Run as: `srun --reserv=openacc -Cgpu ./dot.openacc [ARRAY_SIZE]`
  - `ARRAY_SIZE` is power of 2, default is 2
- Try with different sizes. Does the GPU outperform the CPU version?

# Hands-on

## Exercise 3 – 1D blur kernel



# Hands-on

## Exercise 3 – 1D blur kernel

- `practicals/basics/blur_openacc.{cpp,f90}`
- Run as:  
`srun --reserv=openacc -Cgpu ./blur.openacc [ARRAY_SIZE]`
  - `ARRAY_SIZE` is power of 2, default is 20
- Offload to GPU the loops of the naive kernel; why is it so slow?

# Hands-on

## Exercise 3 – 1D blur kernel

- `practicals/basics/blur_openacc.{cpp,f90}`
- Run as:  
`srun --reserv=openacc -Cgpu ./blur.openacc [ARRAY_SIZE]`
  - `ARRAY_SIZE` is power of 2, default is 20
- Offload to GPU the loops of the naive kernel; why is it so slow?
- Moving data to and from the device is slow ( $\approx 7\text{--}8$  GB/s per direction)
- Avoid unnecessary data movement in the `nocopies` kernel
  - Move the necessary data to GPU early enough and keep it there as long as possible
  - Update host copies using `#pragma acc update` directive if needed

# Hands-on

## Exercise 4 – Experiment with the unified memory

- Remove all the data directives and data clauses
- Compile the `blur_twice_naive` kernel with `-Mcuda=managed`
- How does it compare to the manual data management in terms of performance?
- Can you explain the performance difference?

# Interoperability with CUDA

- Can I use a CUDA pointer inside OpenACC context?
- Can I call a CUDA function from OpenACC context?

Short answer is *yes*.

# Interoperability with CUDA

Use CUDA pointers inside OpenACC context

A scenario:

- Have a CUDA code that needs to call a function that uses OpenACC.
- This function may accept an array that has been allocated already on the GPU by CUDA.

The problem?

- OpenACC only knows of pointers that it is managing itself; the present clause won't work. No idea what this pointer is; never seen it before!

# Interoperability with CUDA

Use CUDA pointers inside OpenACC context

Solution:

- We need to instruct the OpenACC runtime to trust this pointer and that it is a valid device pointer.
- OpenACC runtime will just treat that pointer as known, but it won't check its shape.
- Use the `deviceptr(<ptrlist>)` clause with `parallel`, `kernels` and `data` directives



# Interoperability with CUDA

## Use CUDA pointers inside OpenACC context – Example

```
void copy(double *dst, const double *src, size_t n) {
    #pragma acc parallel loop deviceptr(dst, src)
    for (size_t i = 0; i < n; ++i) {
        dst[i] = src[i];
    }
}

int main() {
    double *a, *b;
    cudaMalloc(&a, 1024);
    cudaMalloc(&b, 1024);
    ...
    copy(b, a, 1024);
    return 0;
}
```

# Interoperability with CUDA

Call a CUDA function from OpenACC context

Scenario:

- My code is in OpenACC, but I need to call an optimized library written in CUDA, which accepts device pointers, e.g., cuBLAS.

# Interoperability with CUDA

Call a CUDA function from OpenACC context

Scenario:

- My code is in OpenACC, but I need to call an optimized library written in CUDA, which accepts device pointers, e.g., cuBLAS.

Problem:

- I only “see” device pointers while in a parallel region, but I want to get a device pointer, while executing on the host.

# Interoperability with CUDA

Call a CUDA function from OpenACC context

Scenario:

- My code is in OpenACC, but I need to call an optimized library written in CUDA, which accepts device pointers, e.g., cuBLAS.

Problem:

- I only “see” device pointers while in a parallel region, but I want to get a device pointer, while executing on the host.

Solution:

- Use a `host_data` region
  - `#pragma acc host_data use_device(<varlist>)`

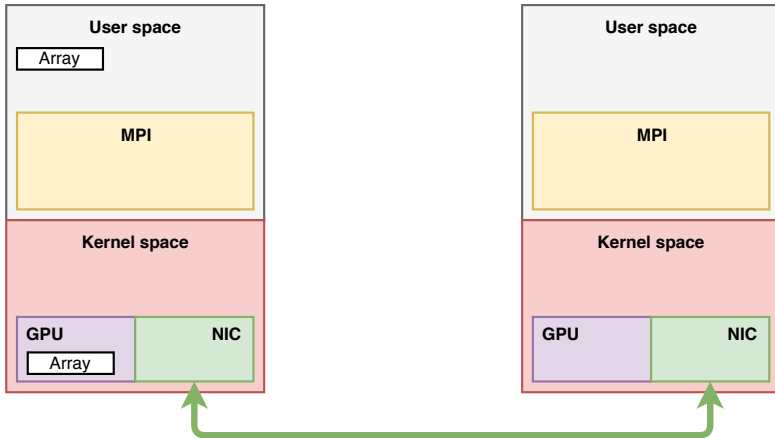
# Interoperability with CUDA

## The `host_data` directive

- C/C++: `#pragma acc host_data use_device(<varlist>)`
  - In the next block of code the compiler will make available the device address of any variable in `<varlist>`.
- Fortran: `!$acc host_data use_device(<varlist>)`
  - The compiler will make available the device address of any variable in `<varlist>` until a matching `!$acc end host_data` is found.
- Optional clauses:
  - `if(condition)`: Use the device pointer if *condition* is true.
  - `if_present`: Use the device pointer if variables in `<varlist>` are present on the device.

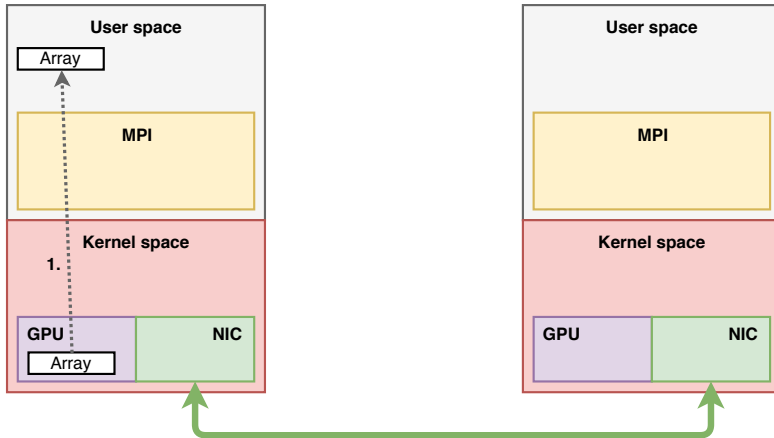
# Interoperability with MPI

The communication data path



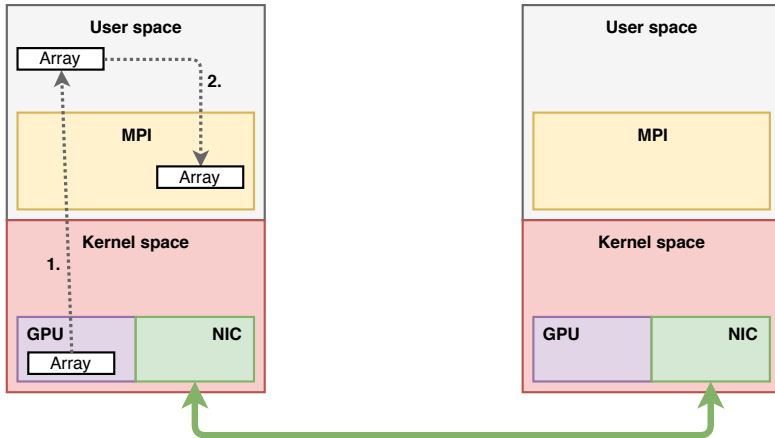
# Interoperability with MPI

The communication data path



# Interoperability with MPI

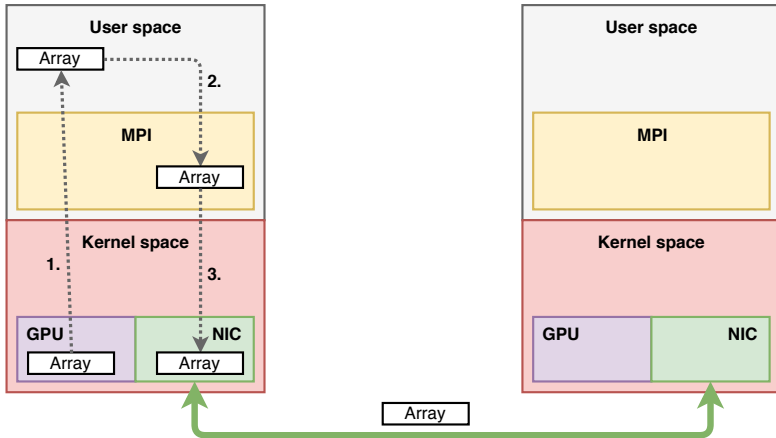
The communication data path





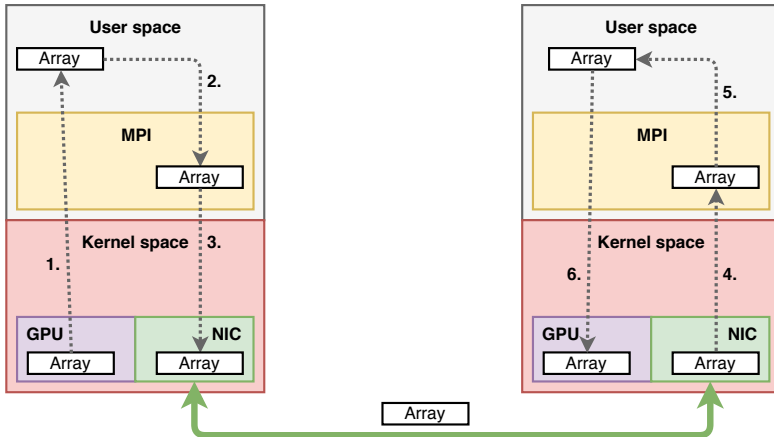
# Interoperability with MPI

The communication data path



# Interoperability with MPI

The communication data path



# Interoperability with MPI

The communication data path

- Aren't there too many copies?
- This path is not fully unoptimized though! It still bypasses a copy by enabling RDMA between the NICs.

# Interoperability with MPI

The communication data path

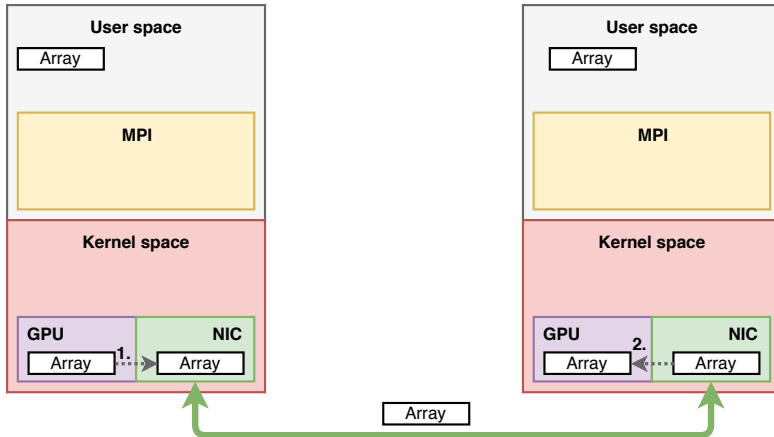
- Aren't there too many copies?
- This path is not fully unoptimized though! It still bypasses a copy by enabling RDMA between the NICs.

Ideally, we would like to avoid all these copies and have the GPU talk directly to the NIC.

- Answer: GPUDirect

# Interoperability with MPI

The GPUDirect optimized RDMA path



# Interoperability with MPI

How to enable this optimized path with OpenACC

- The MPI implementation *must* support it! Cray MPICH implementation does.
  - Enabled by setting `MPICH_RDMA_ENABLED_CUDA=1`.
- If a GPU pointer passed in Send/Recv call, the MPI implementation enables the optimized GPUDirect data path.

# Interoperability with MPI

How to enable this optimized path with OpenACC

- The MPI implementation *must* support it! Cray MPICH implementation does.
  - Enabled by setting `MPICH_RDMA_ENABLED_CUDA=1`.
- If a GPU pointer passed in Send/Recv call, the MPI implementation enables the optimized GPUDirect data path.
- Wrap the Send/Recv calls using the `host_data` directive to get the OpenACC device pointers.

# Hands-on

## Exercise 5 – Calling cuBLAS methods

Source code:

- `practicals/gemm/gemm.cpp`

- Run as:

```
srunk --reserv=openacc -Cgpu ./axpy.openacc [ARRAY_SIZE]
```

- ARRAY\_SIZE is power of 2, default is 16

Steps:

1. Compile with 'make CPPFLAGS=' to get also the naive implementation → too slow!
2. Offload the GEMM method to the GPU using OpenACC
3. Make use of cuBLAS GEMM through OpenACC
4. When does it start to pay of using cuBLAS?



# Hands-on

## Exercise 6.1 – 2D diffusion example

Source code:

- `diffusion2d_omp.{cpp,f90}`: our baseline code
  - Single node OpenMP version for the CPU
- `diffusion2d_openacc.{cpp,f90}`
  - Single node OpenACC version
  - Run as:  
`srun --reserv=openacc -Cgpu ./diffusion2d.openacc [ARRAY_SIZE]`
    - `ARRAY_SIZE` is power of 2, default is 16
  - Fill in the parts where `OPENACC_DATA` is defined.

# Hands-on

## Exercise 6.2 – 2D diffusion example using CUDA data management

Source code:

- `diffusion2d_openacc.{cpp,f90}`

- Single node OpenACC version

- Run as:

- ```
srunk --reserv=openacc -Cgpu ./diffusion2d.openacc.cuda [ARRAY_SIZE]
```

- ARRAY\_SIZE is power of 2, default is 16

- Fill in the parts where OPENACC\_DATA is defined.

# Hands-on

## Exercise 6.3 – 2D diffusion example using MPI

Source code:

- `diffusion2d_openacc_mpi.{cpp,f90}`
  - Run as:  
`srun --reserv=openacc -Cgpu ./diffusion2d.openacc.cuda [ARRAY_SIZE]`
    - `ARRAY_SIZE` is power of 2, default is 16
  - Fill in the parts where `OPENACC_DATA` is defined.